NEURALPATH CHEATSHEET

Module 8: MLOps & Production Systems

Ship reliable models with experiment tracking, ONNX/vLLM serving, quantization, observability,

Core Formulas, Mental Models & Engineering Rules

- Track code SHA, data snapshot hash, hyperparameters, and random seeds per run.
- Serving: ONNX Runtime for classical/vision; vLLM continuous batching + PagedAttention for LLMs.
- Quantization: INT8 / FP8 / AWQ 4-bit cut VRAM 2x-4x with minimal perplexity shift.
- Monitor latency P99, PSI/KS feature drift, and automated golden-suite regression gates.

```
Estimated study time: 4 hours · Badge: Production ML Engineer
```